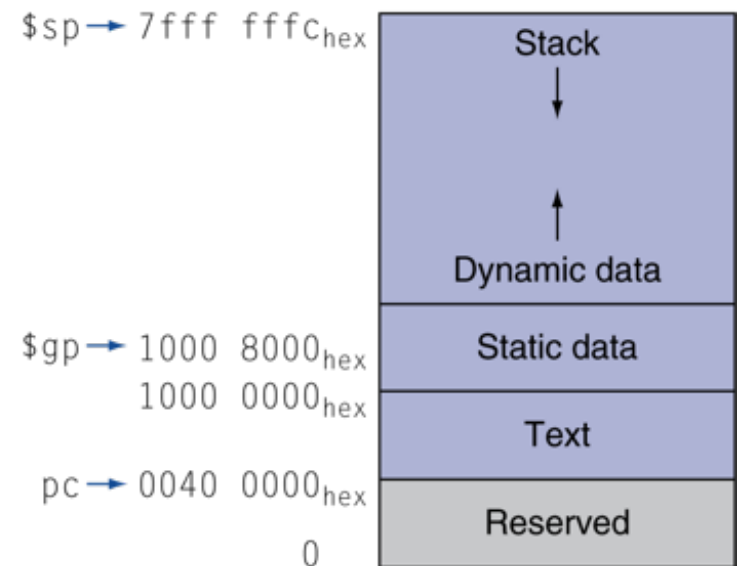


# RoadMap: Procedure Calling

- Steps required
  1. Prepare and pass arguments to callee ✓
    - `$a` registers
  2. Transfer control to callee procedure ✓
    - `jal`
  3. Acquire storage for procedure ✓
    - `Manipulate $sp + lw/sw`
  4. Perform procedure's operations ←
  5. Pass result to caller ✓
    - `$v` registers
  6. Return to place of call ✓
    - `jr $ra`

# Memory Layout

- Text: program code
- Static data: global variables
  - e.g., static variables in C, constant arrays and strings
  - \$gp initialized to address allowing  $\pm$ offsets into this segment
- Dynamic data: heap
  - E.g., malloc in C, new in Java
  - Need a system call in MIPS
- Stack: automatic storage
  - High to low



# Procedure Calling

- Steps required
  1. Prepare and pass arguments to callee ✓
  2. Transfer control to callee procedure ✓
  3. Acquire storage for procedure ✓
  4. Perform procedure's operations
    - How to make procedure calls?
  5. Pass result to caller ✓
  6. Return to place of call ✓

# Non-Leaf Procedure Example

```
main()
```

```
{ ----
```

```
----
```


```
A (3);
```

```
X: ---
```

```
}
```

To call A(3): \$a0 = 3

jal A → \$ra = address X



```
int A (int n)
```

```
{ ---
```



```
B (5);
```

```
Y: ---
```

```
}
```

To call B(5): \$a0 = 5

jal B → \$ra = address Y



```
int B (int m)
```

```
{ ---
```

```
}
```

- Not able to continue to use \$a0 as 3!
- Not able to return to main!

# Non-Leaf Procedures

- Procedures that call other procedures
- Caller needs to prepare for making a call
  - Put parameters in argument registers
  - Set the return address
- If caller is a subroutine, caller needs to save on the stack:
  - Its own return address
  - Any arguments (and temporaries) needed after the call
- Restore from the stack after the call

# Register Conventions

- **Caller** preserved registers
  - Return address: **\$ra** (reg 31)
  - Arguments: **\$a0, \$a1, \$a2, \$a3** (reg's 4 – 7)
  - Return value: **\$v0, \$v1** (reg's 2 and 3)
  - Temporaries: **\$t0, \$t1, ... , \$t7, \$t8, \$t9** (reg's 8 – 15,24,25)
- **Callee** preserved registers
  - Local variables: **\$s0, \$s1, ... , \$s7** (reg's 16 – 23)
  - Stack/frame pointer: **\$sp, \$fp** (reg 29,30)
- Only need to save used registers
- A subroutine can be both caller and callee (non-leaf)

# Non-Leaf Procedure Example

- C code:

```
int fact (int n)
{
    if (n < 1) return 1;
    else return n * fact(n - 1);
}
```

- Argument n in \$a0
- Result in \$v0

# Non-Leaf Procedure Example

fact:

```
addi $sp, $sp, -8      # adjust stack for 2 items
sw    $ra, 4($sp)      # save return address
sw    $a0, 0($sp)      # save argument
```

```
slti  $t0, $a0, 1      # test for n < 1
beq   $t0, $zero, L1
addi  $v0, $zero, 1     # if so, result is 1
```

```
addi  $sp, $sp, 8       # pop 2 items from stack
jr    $ra               # and return
```

```
L1: addi $a0, $a0, -1    # else decrement n
jal   fact              # recursive call
```

```
lw    $a0, 0($sp)       # restore original n
lw    $ra, 4($sp)       # and return address
addi  $sp, $sp, 8       # pop 2 items from stack
```

```
mul   $v0, $a0, $v0     # multiply to get result
jr    $ra               # and return
```

```
int fact (int n)
{
    if (n < 1) return 1;
    else return n * fact(n - 1);
}
```



# Steps for Making a Procedure Call

- 1) Save necessary values into stack
  - Return address, arguments, ...
- 2) Assign argument(s), if any
  - Special registers
- 3) jal call
  - Control transferred to callee after executing this instruction
- 4) Restore values from stack after return from callee

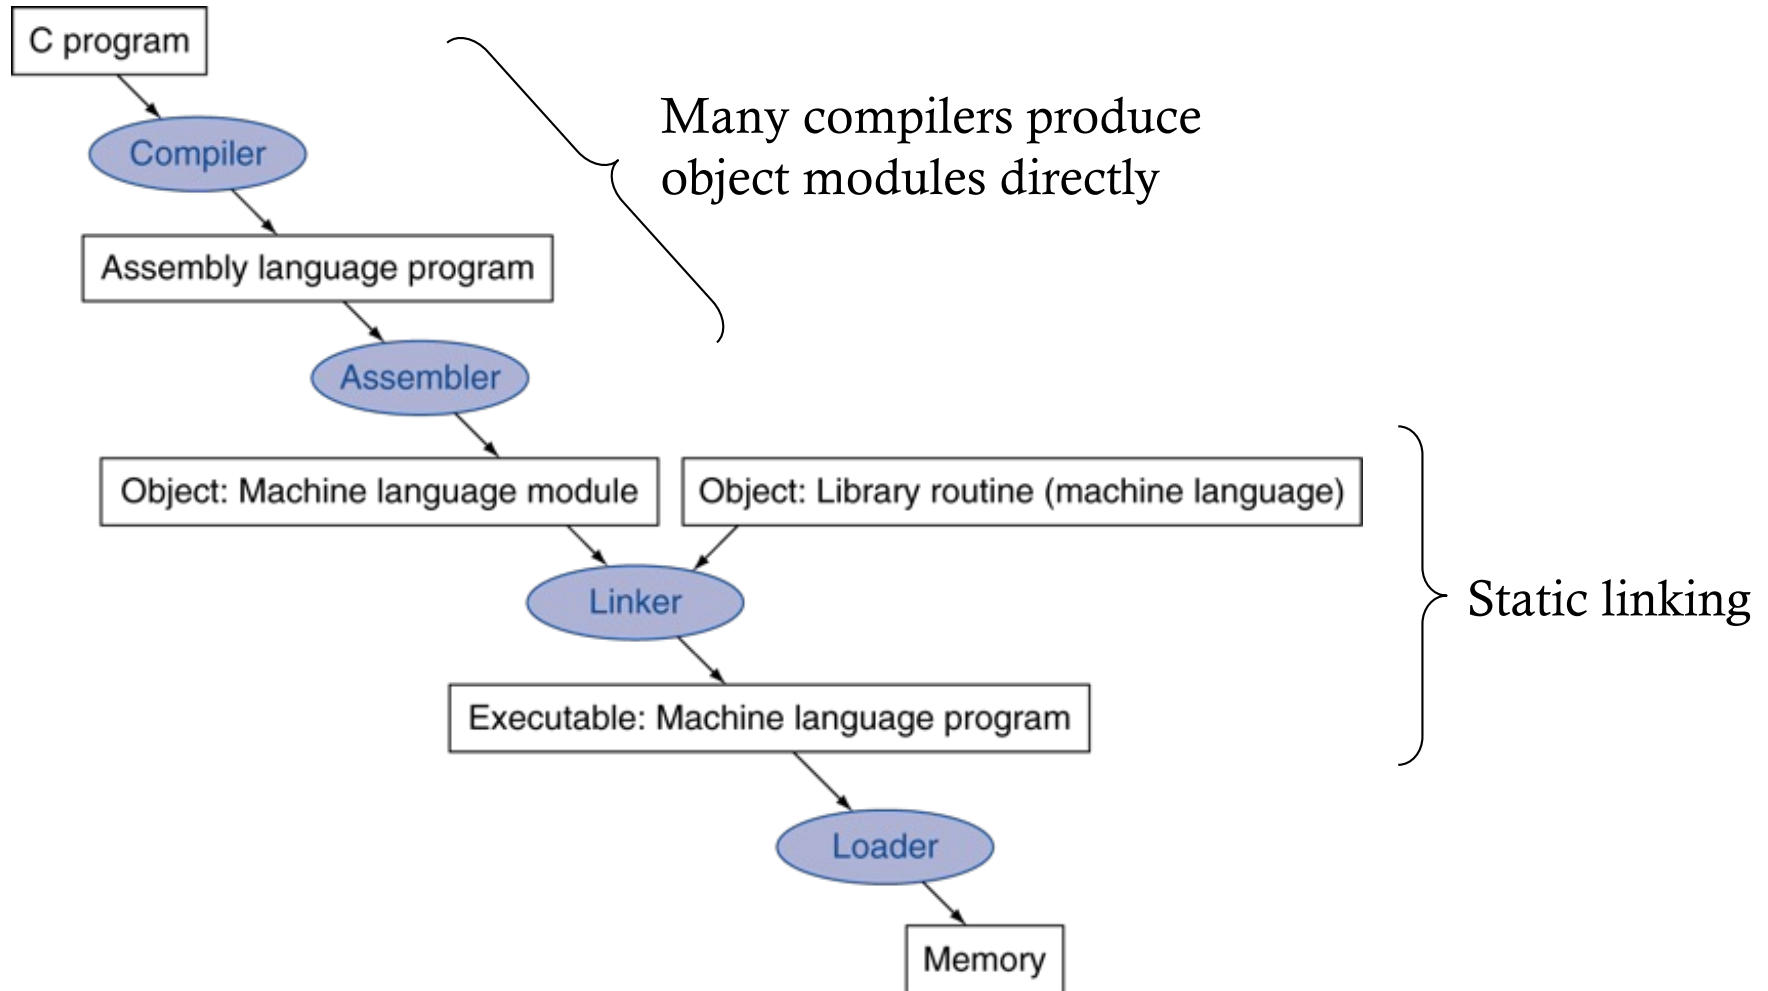
**Prolog**

**Epilog**

# Road Map – MIPS ISA

- MIPS basic instructions
- MIPS instruction format
- Procedure calls
- Other misc.
  - Translating and starting a program
  - Synchronization instructions
  - Arrays and pointers
  - Fallacies and pitfalls

# Translation and Startup



# Assembler Pseudo-instructions

- Most assembler instructions represent machine instructions one-to-one
- Pseudo-instructions: figments of the assembler's imagination

`move $t0, $t1`       $\rightarrow$  `add $t0, $zero, $t1`

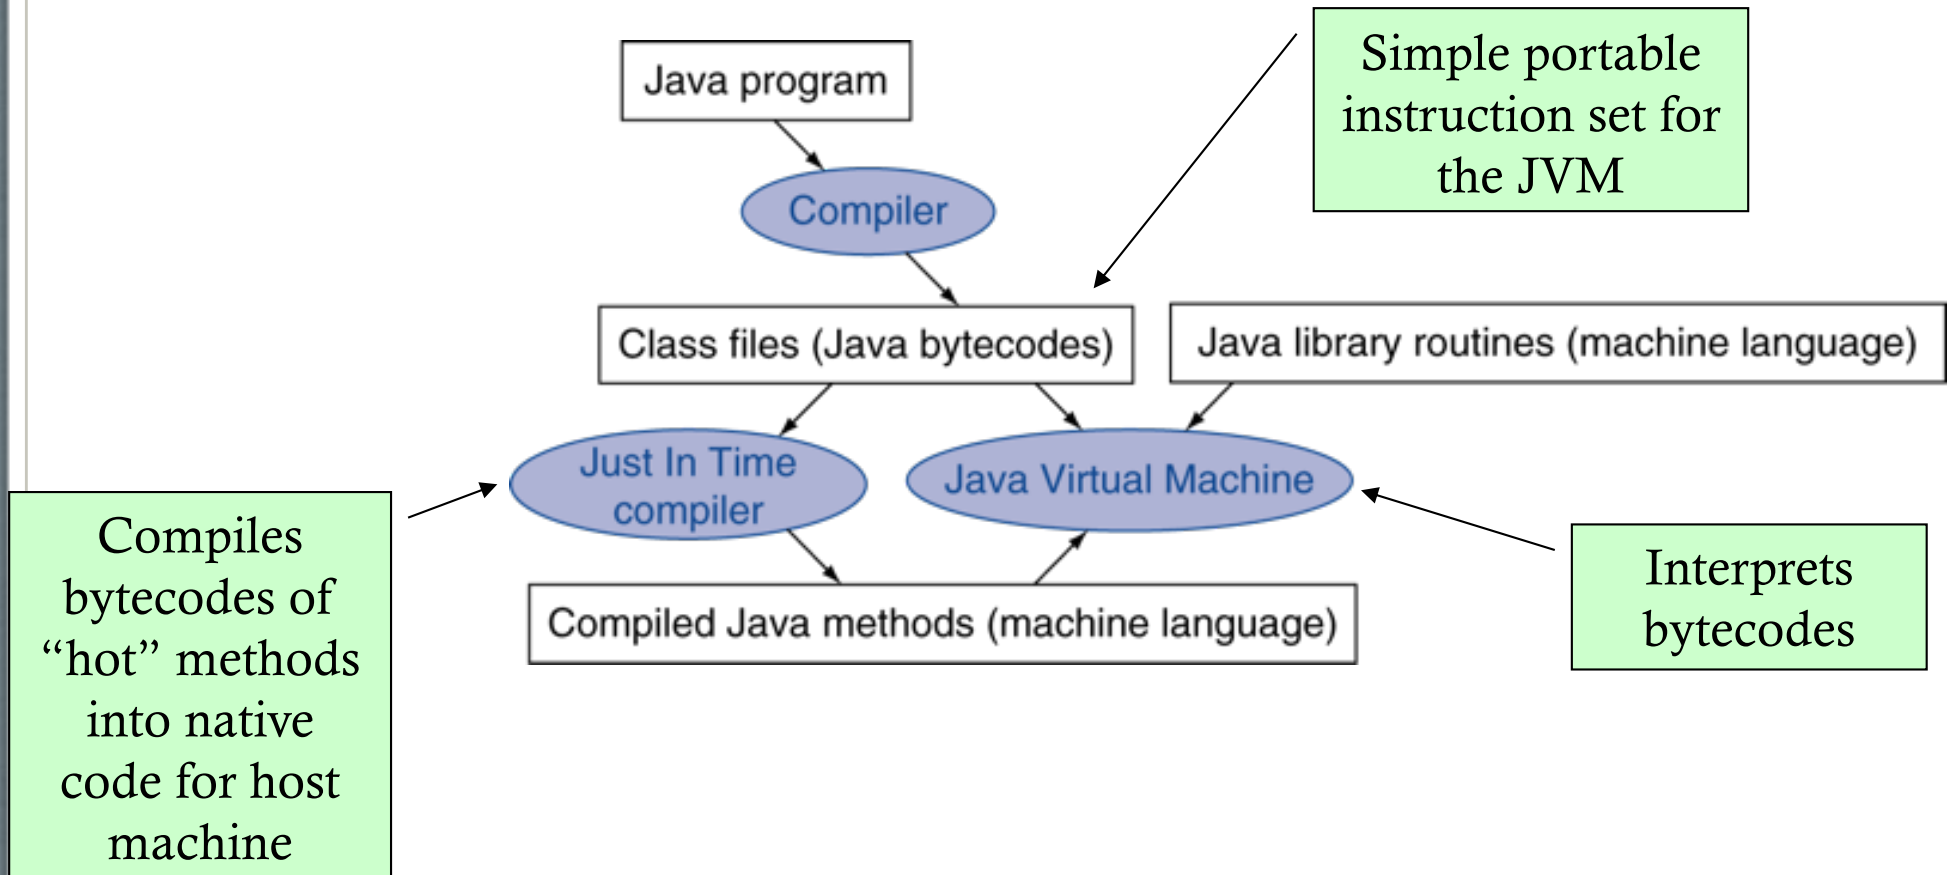
`blt $t0, $t1, L`       $\rightarrow$  `slt $at, $t0, $t1`  
                             `bne $at, $zero, L`

- `$at` (register 1): assembler temporary

# Label Resolution

- Labels need to be translated into addresses
  - Branch/jump target, data addresses
- Assembler (or compiler) translates program into machine instructions
  - Label → binary value if possible
  - Symbol table: remaining labels that are undefined in this module
- Linker uses relocation information and symbol table in each object module to resolve undefined labels
  - DLL may delay library linking to runtime

# Starting Java Applications



# Effect of Language, Compiler, and Algorithm




Fig. 2.29

# Example: Data Races

Processor P1	Processor P2
1.1 Read X from memory	2.1 Read X from memory
1.2 Add 1 to X (local copy)	2.2 Add 2 to X (local copy)
1.3 Write X to memory	2.3 Write X to memory

Possible sequence: **All updates are applied!**

Operation	P1	P2	Memory
			X = 100
1.1	X = 100		X = 100
1.2	X = 101		X = 100
1.3	X = 101		X = 101
2.1		X = 101	X = 101
2.2		X = 103	X = 101
2.3		X = 103	X = 103



# Example: Data Races

Processor P1	Processor P2
1.1 Read X from memory	2.1 Read X from memory
1.2 Add 1 to X (local copy)	2.2 Add 2 to X (local copy)
1.3 Write X to memory	2.3 Write X to memory

Possible sequence: **Inconsistent / lost updates!**

Operation	P1	P2	Memory
			X = 100
1.1	X = 100		X = 100
2.1	X = 100	X = 100	X = 100
1.2	X = 101	X = 100	X = 100
2.2	X = 101	X = 102	X = 100
2.3	X = 101	X = 102	X = 102
1.3	X = 101	X = 102	X = 101

# Synchronization

- Parallel computing: coordination needed between multiple cooperating tasks
  - Enforce a certain order / avoid data races
- Hardware support required to ensure atomic read/modify memory operations
  - Atomic: no other access to the location allowed between the read and write
  - Basis of building locks / critical sections / ...
- Implementation options
  - A single atomic memory instruction
  - A pair of instructions (more efficient, may succeed/fail)

# Synchronization in MIPS

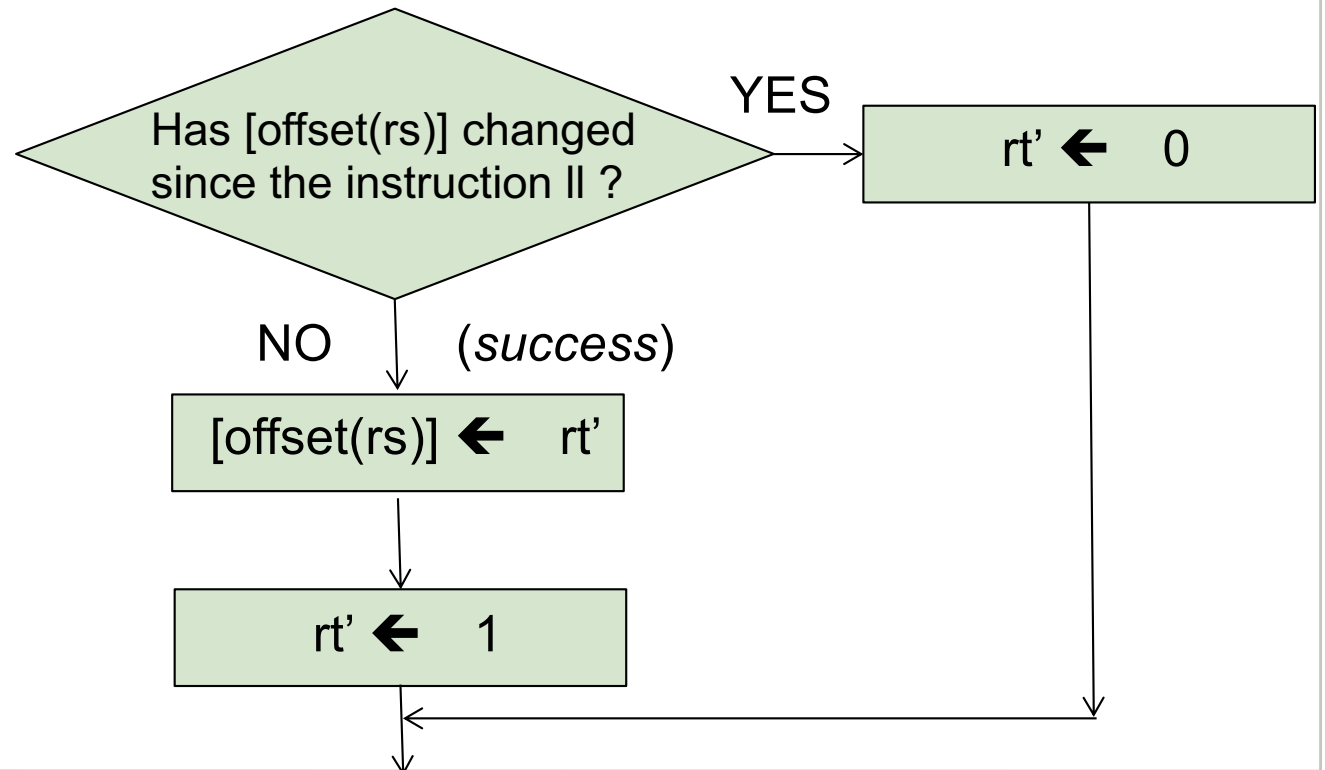
- Load linked: `ll rt, offset(rs)`
- Store conditional: `sc rt, offset(rs)`
  - Always set `rt` as a report
  - Return 1 in `rt`
    - Succeed/atomic: memory location not changed since the previous ll
  - Return 0 in `rt`
    - Fail: if location has been changed since the previous ll

# Synchronization in MIPS

II  $rt, \text{offset}(rs)$

$rt \leftarrow [\text{offset}(rs)]$  (start monitoring  $[\text{offset}(rs)]$ )

sc  $rt', \text{offset}(rs)$



# Uses of LL-SC

- Atomic updating one memory location
  - E.g.  $a[0] = a[0] + 1$

```
try: ll    $t0,offset($s1)    #load linked
      addi  $t0,$t0,1          #increment by 1
      sc    $t0,offset($s1)    #store conditional
      beq   $t0,$zero,try       #re-try if fails

      << successfully updated offset($s1) >>
```

- Locks, critical sections, ...

# Arrays vs. Pointers

- Array indexing involves
  - Multiplying index by element size
  - Adding to array base address
- Pointers correspond directly to memory addresses
  - Can avoid indexing complexity

# Example: Clearing an Array

```
clear1(int array[], int size) {
    int i;
    for (i = 0; i < size; i += 1)
        array[i] = 0;
}
```

```
        move $t0,$zero    # i = 0
        j test1
loop1:  sll $t1,$t0,2      # $t1 = i * 4
        add $t2,$a0,$t1   # $t2 =
                                # &array[i]
        sw $zero, 0($t2)  # array[i] = 0
        addi $t0,$t0,1    # i = i + 1
test1:  slt $t3,$t0,$a1    # $t3 =
                                # (i < size)
        bne $t3,$zero,loop1 # if (...)
                                # goto loop1
```

```
clear2(int *array, int size) {
    int *p;
    for (p = &array[0]; p < &array[size];
        p = p + 1)
        *p = 0;
}
```

```
        move $t0,$a0      # p = & array[0]
        sll $t1,$a1,2      # $t1 = size * 4
        add $t2,$a0,$t1    # $t2 =
                                # &array[size]
        j test2
loop2:  sw $zero,0($t0)    # Memory[p] = 0
        addi $t0,$t0,4     # p = p + 4
test2:  slt $t3,$t0,$t2    # $t3 =
                                # (p < &array[size])
        bne $t3,$zero,loop2 # if (...)
                                # goto loop2
```

\*code changed slightly from textbook 2.14

# Pitfalls

- Sequential words are not at sequential addresses
  - Increment by 4, not by 1!
- Keeping a pointer to an automatic variable after procedure returns
  - e.g., passing pointer back via an argument
  - Pointer becomes invalid when stack popped




# Fallacies

- Powerful instruction  $\Rightarrow$  higher performance
  - Fewer instructions required
  - But complex instructions are hard to implement
    - May slow down all instructions, including simple ones
  - Compilers are good at making fast code from simple instructions
- Assembly programs: manually developed vs. automatically generated
  - Modern compilers improves a lot for performance
  - More lines of code  $\Rightarrow$  more errors and less productivity

# Fallacies

- Backward compatibility  $\Rightarrow$  instruction set doesn't change
  - But they do accrete more instructions



x86 instruction set

# MIPS Common Instructions

- Measure MIPS instruction executions in benchmark programs
  - Consider making the common case fast

Instruction class	MIPS examples	SPEC2006 Int	SPEC2006 FP
Arithmetic	add, sub, addi	16%	48%
Data transfer	lw, sw, lb, lbu, lh, lhu, sb, lui	35%	36%
Logical	and, or, nor, andi, ori, sll, srl	12%	4%
Cond. Branch	beq, bne, slt, slti, sltiu	34%	8%
Jump	j, jr, jal	2%	0%

# Concluding Remarks

- Design principles
  1. Simplicity favors regularity
  2. Smaller is faster
  3. Make the common case fast
  4. Good design demands good compromises
- Layers of software/hardware
  - Compiler, assembler, hardware
- MIPS: typical of RISC ISAs
  - c.f. x86