

Roadmap

- Resolving hazards
 - Data hazards: forwarding
 - Stalling the pipeline
 - Control hazards: prediction
- Advanced ILP techniques
 - Static approaches
 - Reordering, multiple issue, **loop unrolling**, ...
 - Dynamic scheduling and speculation
- Concluding remarks

Scheduling Example

- Schedule this for dual-issue MIPS

```

Loop: lw    $t0, 0($s1)      # $t0=array element
      addu  $t0, $t0, $s2    # add scalar in $s2
      sw    $t0, 0($s1)     # store result
      addi  $s1, $s1, -4     # decrement pointer
      bne   $s1, $zero, Loop # branch $s1!=0
    
```

	ALU/branch	Load/store	cycle
Loop:	nop	lw \$t0, 0(\$s1)	1
	addi \$s1, \$s1, -4	nop	2
	addu \$t0, \$t0, \$s2	nop	3
	bne \$s1, \$zero, Loop	sw \$t0, 4(\$s1)	4

- $IPC = 5/4 = 1.25$ (c.f. peak $IPC = 2$)

Loop Unrolling

- Replicate loop body and adjust the increment of loop variable and loop body
 - Use different registers per replication - “register renaming”
- Effects
 - Loop overhead is reduced
 - Larger loop body: more ILP/ scheduling freedom
 - Register usage within the loop body increased

do i=1 to n by 1		do i=1 to n by 2
a[i] = a[i]*s	⇒	a[i] = a[i]*s
end		a[i+1] = a[i+1]*s
		end

Additional code needed at the end if n is not even

Loop Unrolling Example

- Original loop

```
Loop: lw    $t0, 0($s1)      # $t0=array element
      addu  $t0, $t0, $s2    # add scalar in $s2
      sw    $t0, 0($s1)      # store result
      addi  $s1, $s1, -4     # decrement pointer
      bne   $s1, $zero, Loop # branch $s1!=0
```

- I: make four copies of the loop body

Loop Unrolling Example

Loop: lw \$t0, 0(\$s1) # \$t0=array element
 addu \$t0, \$t0, \$s2 # add scalar in \$s2
 sw \$t0, 0(\$s1) # store result
 addi \$s1, \$s1, -4 # decrement pointer
 bne \$s1, \$zero, Loop # branch \$s1!=0

 lw \$t0, 0(\$s1) # \$t0=array element
 addu \$t0, \$t0, \$s2 # add scalar in \$s2
 sw \$t0, 0(\$s1) # store result
 addi \$s1, \$s1, -4 # decrement pointer
 bne \$s1, \$zero, Loop # branch \$s1!=0

 lw \$t0, 0(\$s1) # \$t0=array element
 addu \$t0, \$t0, \$s2 # add scalar in \$s2
 sw \$t0, 0(\$s1) # store result
 addi \$s1, \$s1, -4 # decrement pointer
 bne \$s1, \$zero, Loop # branch \$s1!=0

 lw \$t0, 0(\$s1) # \$t0=array element
 addu \$t0, \$t0, \$s2 # add scalar in \$s2
 sw \$t0, 0(\$s1) # store result
 addi \$s1, \$s1, -4 # decrement pointer
 bne \$s1, \$zero, Loop # branch \$s1!=0

- 2: drop unnecessary updating and checking of loop variable – only need to keep one checking/updating

Loop Unrolling Example

```
Loop:  lw  $t0, 0($s1)    # $t0=array element
      addu $t0, $t0, $s2  # add scalar in $s2
      sw  $t0, 0($s1)    # store result

      lw  $t0, 0($s1)    # $t0=array element
      addu $t0, $t0, $s2  # add scalar in $s2
      sw  $t0, 0($s1)    # store result

      lw  $t0, 0($s1)    # $t0=array element
      addu $t0, $t0, $s2  # add scalar in $s2
      sw  $t0, 0($s1)    # store result

      lw  $t0, 0($s1)    # $t0=array element
      addu $t0, $t0, $s2  # add scalar in $s2
      sw  $t0, 0($s1)    # store result
      addi $s1, $s1, -16  # decrement pointer
      bne $s1, $zero, Loop # branch $s1!=0
```

- 3: offsets needs to be adjusted for different iterations

Loop Unrolling Example

Loop: lw \$t0, 0(\$s1) # \$t0=array element
 addu \$t0, \$t0, \$s2 # add scalar in \$s2
 sw \$t0, 0(\$s1) # store result

 lw \$t0, -4(\$s1) # \$t0=array element
 addu \$t0, \$t0, \$s2 # add scalar in \$s2
 sw \$t0, -4(\$s1) # store result

Dependences?

 lw \$t0, -8(\$s1) # \$t0=array element
 addu \$t0, \$t0, \$s2 # add scalar in \$s2
 sw \$t0, -8(\$s1) # store result

 lw \$t0, -12(\$s1) # \$t0=array element
 addu \$t0, \$t0, \$s2 # add scalar in \$s2
 sw \$t0, -12(\$s1) # store result
 addi \$s1, \$s1, -16 # decrement pointer
 bne \$s1, \$zero, Loop # branch \$s1!=0

- 4: rename registers to remove anti- and output dependences

Loop Unrolling Example

Loop: lw \$t0, 0(\$s1) # \$t0=array element
 addu \$t0, \$t0, \$s2 # add scalar in \$s2
 sw \$t0, 0(\$s1) # store result

 lw \$t1, -4(\$s1) # \$t1=array element
 addu \$t1, \$t1, 0 \$s2 # add scalar in \$s2
 sw \$t1, -4(\$s1) # store result

 lw \$t2, -8(\$s1) # \$t2=array element
 addu \$t2, \$t2, \$s2 # add scalar in \$s2
 sw \$t2, -8(\$s1) # store result

 lw \$t3, -12(\$s1) # \$t3=array element
 addu \$t3, \$t3, \$s2 # add scalar in \$s2
 sw \$t3, -12(\$s1) # store result
 addi \$s1, \$s1, -16 # decrement pointer
 bne \$s1, \$zero, Loop # branch \$s1!=0

- 5: ready to schedule!

Scheduling Unrolled Loop

	ALU/branch	Load/store	cycle
Loop:	addi \$s1 , \$s1, -16	lw \$t0 , 0(\$s1)	1
	nop	lw \$t1 , 12(\$s1)	2
	addu \$t0 , \$t0 , \$s2	lw \$t2 , 8(\$s1)	3
	addu \$t1 , \$t1 , \$s2	lw \$t3 , 4(\$s1)	4
	addu \$t2 , \$t2 , \$s2	sw \$t0 , 16(\$s1)	5
	addu \$t3 , \$t4 , \$s2	sw \$t1 , 12(\$s1)	6
	nop	sw \$t2 , 8(\$s1)	7
	bne \$s1 , \$zero, Loop	sw \$t3 , 4(\$s1)	8


- $IPC = 14/8 = 1.75$
 - Closer to 2, but at cost of registers and code size

Outline

- Resolving hazards
 - Data hazards: forwarding
 - Stalling the pipeline
 - Control hazards: prediction
- Advanced ILP techniques
 - Static approaches
 - Reordering, multiple issue, loop unrolling
 - Dynamic scheduling and speculation
- Concluding remarks

Dynamic Scheduling

- Key idea: determine which instructions are ready and allow independent instructions behind stall to proceed




Better if we can start sub while addu is waiting for lw – How?

- Simple pipeline
 - Issue stage: ID (instruction decode)
 - Decoding / check hazards
 - Read operands

Dynamic Scheduling

- Split the ID pipe stage of simple 5-stage pipeline into 2 stages:
 - *Issue* — decode instructions, check for hazards (no delay needed)
 - *Read operands* — wait until no data hazards, then read operands (execution can start immediately after operands reading)

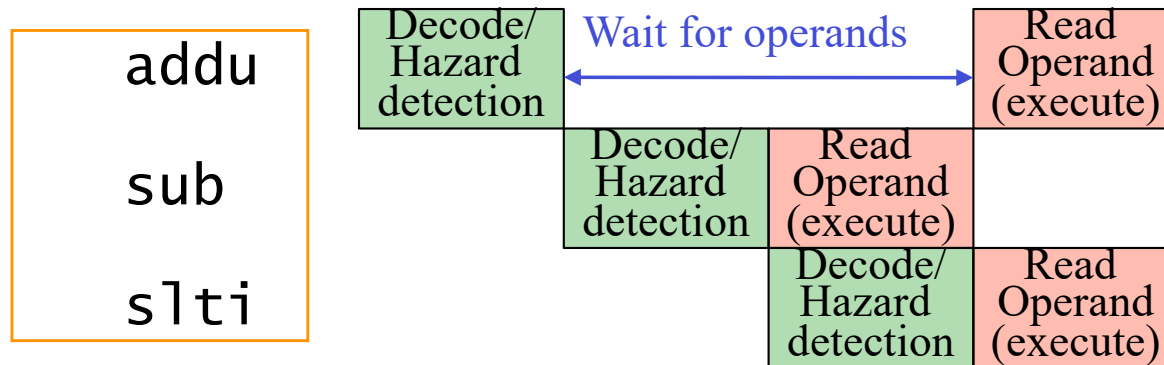


- Enables *in-order issue*, *out-of-order execution* and allows *out-of-order completion*

Splitting Issue & Execution

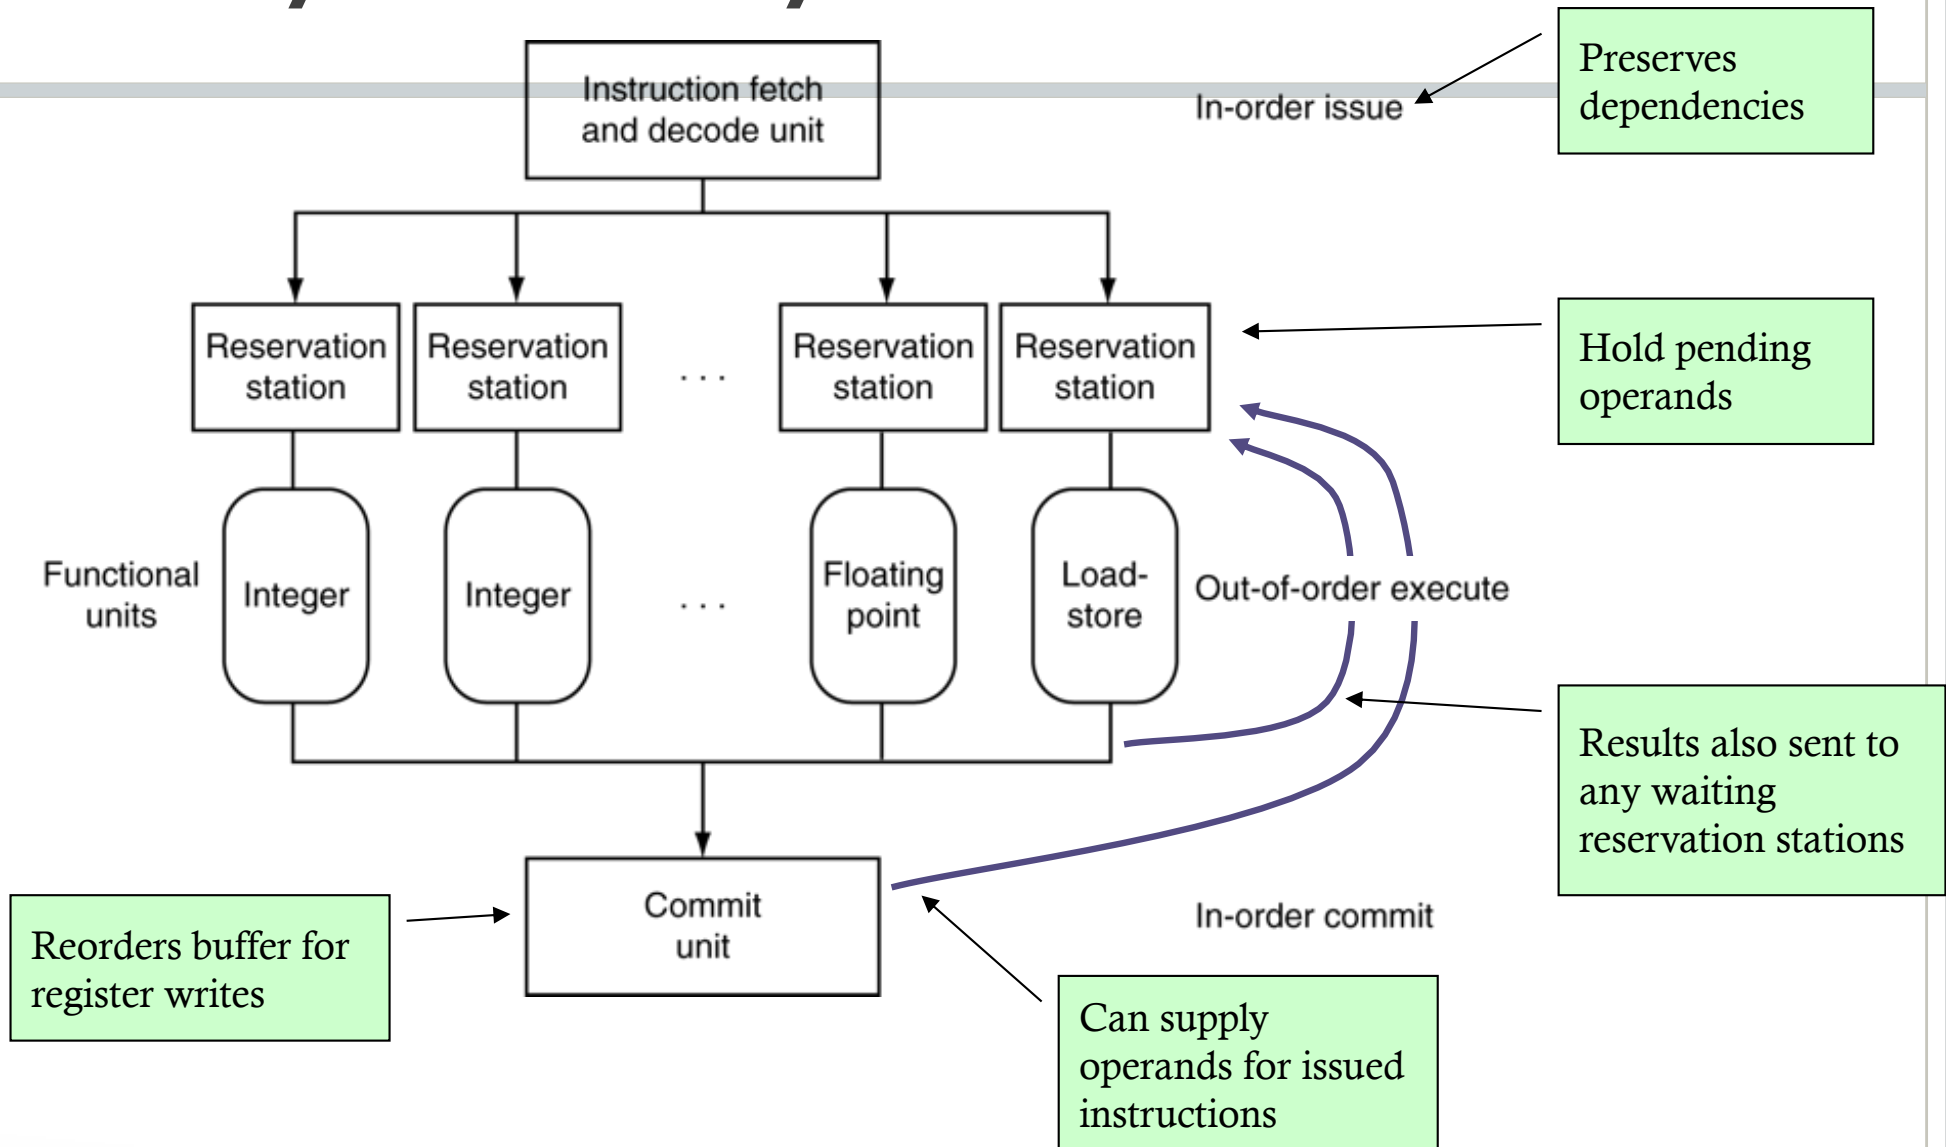
```
lw    $t0, 20($s2)
addu  $t1, $t0, $t2
sub    $s4, $s4, $t3
slti   $t5, $s4, 20
```

- Need a buffer for instructions which are waiting for the operand(s)



- Need to make sure to forward values to instructions waiting for them

Dynamically Scheduled CPU



Optional: Tomasulo Algorithm

- Hardware implementation of dynamic scheduling
 - First proposed by Tomasulo with IBM
 - Adopted by many processors including Intel x86
- http://en.wikipedia.org/wiki/Tomasulo_algorithm
- <http://courses.cs.washington.edu/courses/csep548/06au/lectures/tomasulo.pdf>

Why Dynamic Scheduling?

- Why not just let the compiler schedule code?
- Not all stalls are predictable
 - e.g., cache misses
- Can't always schedule around branches
 - Branch outcome is dynamically determined
- Different implementations of an ISA have different latencies and hazards

Speculation for Greater ILP

- “Guess” what to do with an instruction
 - Start operation as soon as possible
 - Check whether guess was right
 - If so, complete the operation
 - If not, roll-back and do the right thing
- Common to static and dynamic multiple issue
- Examples
 - Speculate on branch outcome
 - Roll back if path taken is different
 - Speculate on load
 - Roll back if location is updated

Meltdown and Spectre

- Vulnerabilities in modern computers
 - Reported Jan. 2018
 - Affect a wide range of systems
 - Intel x86, IBM POWER, ARM
 - iOS, Linux, macOS, Windows
 - Attacking processors that support dynamic scheduling / speculated execution



Example

```
t = a+b
u = t+c
v = u+d
if v:
    w = kern_mem[address]
    # if we get here, fault
    x = w&0x100
    y = user_mem[x]
```

Original sequence

```
#speculative execution attempted
t, w_ = a+b, kern_mem[address]
u, x_ = t+c, w_&0x100
v, y_ = u+d, user_mem[x_]
if v:
    # fault
    w, x, y = w_, x_, y_ # we never get here
```

Possible shuffled sequence for an out-of-order, dual-issue, speculative processor

- If we think the branch might be taken

Example

- Feels safe
 - Either v is zero,
 - Or kernel access fault
- Race condition in reality
 - Memory accessed and cached before privilege checking completed

```
t, w_ = a+b, kern_mem[address]
u, x_ = t+c, w_ & 0x100
v, y_ = u+d, user_mem[x_]
if v:
    # fault
    w, x, y = w_, x_, y_ # we never get here
```

- Side-channel / timing attack
 - Check cache → get to know which location has been fetched from memory for y_
 - → get the value of x_
 - → deduct one bit of w_, which is from kernel space!

References

- [http://en.wikipedia.org/wiki/Meltdown_\(security_vulnerability\)](http://en.wikipedia.org/wiki/Meltdown_(security_vulnerability))
- [http://en.wikipedia.org/wiki/Spectre_\(security_vulnerability\)](http://en.wikipedia.org/wiki/Spectre_(security_vulnerability))
- <http://meltdownattack.com/>
- <http://googleprojectzero.blogspot.com/2018/01/reading-privileged-memory-with-side.html>
- <http://www.raspberrypi.org/blog/why-raspberry-pi-isnt-vulnerable-to-spectre-or-meltdown/>

Does Advanced ILP Work?

The BIG Picture

- Helped a lot but now comes with diminishing returns
- Programs have real dependencies that limit ILP
- Some dependencies are hard to eliminate
 - e.g., pointer aliasing
- Some parallelism is hard to expose
 - Limited window size during instruction issue
- Memory delays and limited bandwidth
 - Hard to keep pipelines full
- Speculation can help if done well

Power Efficiency

- Complexity of dynamic scheduling and speculations requires power
- Multiple simpler cores may be better

Microprocessor	Year	Clock Rate	Pipeline Stages	Issue width	Out-of-order/ Speculation	Cores	Power
i486	1989	25MHz	5	1	No	1	5W
Pentium	1993	66MHz	5	2	No	1	10W
Pentium Pro	1997	200MHz	10	3	Yes	1	29W
P4 Willamette	2001	2000MHz	22	3	Yes	1	75W
P4 Prescott	2004	3600MHz	31	3	Yes	1	103W
Core	2006	2930MHz	14	4	Yes	2	75W
UltraSparc III	2003	1950MHz	14	4	No	1	90W
UltraSparc T1	2005	1200MHz	6	1	No	8	70W

Concluding Remarks

- ISA influences design of datapath and control
- Datapath and control influence design of ISA
- Pipelining improves instruction throughput using parallelism
 - More instructions completed per second
 - Latency for each instruction not reduced
- Hazards: structural, data, control
- Multiple issue and dynamic scheduling (ILP)
 - Dependencies limit achievable parallelism
 - Complexity leads to the power wall